

Experiment No. 02

Semester	B.E. Semester VII – Computer Engineering
Subject	Big Data Analysis
Subject Professor In-charge	Prof. Pankaj Vanvari
Lab Professor In-charge	Dr. Umesh Kulkarni
Academic Year	2024-25
Student Name	Deep Salunkhe
Roll Number	21102A0014

Title: Create directories and files in HDFS and perform read/write operations

Output:

```
deepsalunkhee@LAPTOP-UGI  ×  +  ∨


deepsalunkhee@LAPTOP-UGU6DF82:~$ start-dfs.sh
Starting namenodes on [localhost]
Starting datanodes
Starting secondary namenodes [LAPTOP-UGU6DF82]
deepsalunkhee@LAPTOP-UGU6DF82:~$ start-yarn.sh
Starting resourcemanager
Starting nodemanagers
deepsalunkhee@LAPTOP-UGU6DF82:~$

deepsalunkhee@LAPTOP-UGU6DF82:~$ jps
8112 SecondaryNameNode
7762 NameNode
8739 ResourceManager
8862 NodeManager
9326 Jps
deepsalunkhee@LAPTOP-UGU6DF82:~$
```

localhost:9870/logs/

Directory: /logs/

Name ↑	Last Modified	Size
hadoop-deepsalunkhee-datanode-LAPTOP-UGU6DF82.log	Sep 30, 2024 7:31:01 PM	256,257 bytes
hadoop-deepsalunkhee-datanode-LAPTOP-UGU6DF82.out	Sep 30, 2024 7:30:58 PM	819 bytes
hadoop-deepsalunkhee-datanode-LAPTOP-UGU6DF82.out.1	Sep 30, 2024 7:24:10 PM	819 bytes
hadoop-deepsalunkhee-datanode-LAPTOP-UGU6DF82.out.2	Sep 30, 2024 7:19:07 PM	819 bytes
hadoop-deepsalunkhee-datanode-LAPTOP-UGU6DF82.out.3	Jul 21, 2024 9:15:35 PM	819 bytes
hadoop-deepsalunkhee-namenode-LAPTOP-UGU6DF82.log	Sep 30, 2024 7:37:00 PM	264,367 bytes
hadoop-deepsalunkhee-namenode-LAPTOP-UGU6DF82.out	Sep 30, 2024 7:30:56 PM	819 bytes
hadoop-deepsalunkhee-namenode-LAPTOP-UGU6DF82.out.1	Sep 30, 2024 7:28:50 PM	819 bytes
hadoop-deepsalunkhee-namenode-LAPTOP-UGU6DF82.out.2	Sep 30, 2024 7:24:08 PM	819 bytes
hadoop-deepsalunkhee-namenode-LAPTOP-UGU6DF82.out.3	Sep 30, 2024 7:19:05 PM	819 bytes
hadoop-deepsalunkhee-namenode-LAPTOP-UGU6DF82.out.4	Jul 21, 2024 9:22:14 PM	6,674 bytes
hadoop-deepsalunkhee-nodemanager-LAPTOP-UGU6DF82.log	Sep 30, 2024 7:32:49 PM	151,274 bytes
hadoop-deepsalunkhee-nodemanager-LAPTOP-UGU6DF82.out	Sep 30, 2024 7:32:48 PM	2,330 bytes
hadoop-deepsalunkhee-nodemanager-LAPTOP-UGU6DF82.out.1	Sep 30, 2024 7:19:31 PM	2,330 bytes
hadoop-deepsalunkhee-nodemanager-LAPTOP-UGU6DF82.out.2	Jul 21, 2024 9:15:47 PM	2,330 bytes
hadoop-deepsalunkhee-resourcemanager-LAPTOP-UGU6DF82.log	Sep 30, 2024 7:32:49 PM	178,493 bytes
hadoop-deepsalunkhee-resourcemanager-LAPTOP-UGU6DF82.out	Sep 30, 2024 7:32:46 PM	2,344 bytes
hadoop-deepsalunkhee-resourcemanager-LAPTOP-UGU6DF82.out.1	Sep 30, 2024 7:19:30 PM	2,344 bytes
hadoop-deepsalunkhee-resourcemanager-LAPTOP-UGU6DF82.out.2	Jul 21, 2024 9:15:45 PM	2,346 bytes


NE

Cluster

About
Nodes
Node Labels
Applications
NEW
NEW_SAVING
SUBMITTED
ACCEPTED
RUNNING
FINISHED
FAILED
KILLED
Scheduler

Tools

Cluster Metrics

Apps Submitted	Apps Pending	Apps Running	Apps Completed	Containers Running	Used
0	0	0	0	0	<memory:0 B, vC

Cluster Nodes Metrics

Active Nodes	Decommissioning Nodes	Decommissioned Nodes
1	0	0

Scheduler Metrics

Scheduler Type	Scheduling Resource Type	Minimum Allocation	Maximum Allocation
Capacity Scheduler	[memory-mb (unit=M), vcores]	<memory:1024, vCores:1>	<memory:8192, vCores:4>

Show 20 entries

ID	User	Name	Application Type	Application Tags	Queue	Application Priority	StartTime	LaunchTime	FinishTime	State	Fi
Showing 0 to 0 of 0 entries											

```

deepsalunkhee@LAPTOP-UGI x + v
deepsalunkhee@LAPTOP-UGU6DF82:~/docs$ cat temp.txt
This is the file from docs folder
deepsalunkhee@LAPTOP-UGU6DF82:~/docs$ cd ..
deepsalunkhee@LAPTOP-UGU6DF82:~$ cd docs
deepsalunkhee@LAPTOP-UGU6DF82:~/docs$ pwd
/home/deepsalunkhee/docs
deepsalunkhee@LAPTOP-UGU6DF82:~/docs$ cd ..
deepsalunkhee@LAPTOP-UGU6DF82:~$ $ hdfs dfs -put /home/deepsalunkhee/docs/temp.txt /user/deep/hadoop-dir
$: command not found
deepsalunkhee@LAPTOP-UGU6DF82:~$ hdfs dfs -put /home/deepsalunkhee/docs/temp.txt /user/deep/hadoop-dir
2024-09-30 19:47:04,120 WARN hdfs.DataStreamer: DataStreamer Exception
org.apache.hadoop.ipc.RemoteException(java.io.IOException): File /user/deep/hadoop-dir/temp.txt_COPYING_ could only be written to 0
of the 1 minReplication nodes. There are 0 datanode(s) running and 0 node(s) are excluded in this operation.
    at org.apache.hadoop.hdfs.server.blockmanagement.BlockManager.chooseTarget4NewBlock(BlockManager.java:2473)
    at org.apache.hadoop.hdfs.server.namenode.FSDirWriteFileOp.chooseTargetForNewBlock(FSDirWriteFileOp.java:293)
    at org.apache.hadoop.hdfs.server.namenode.FSNamesystem.getAdditionalBlock(FSNamesystem.java:3075)
    at org.apache.hadoop.hdfs.server.namenode.NameNodeRpcServer.addBlock(NameNodeRpcServer.java:932)
    at org.apache.hadoop.hdfs.protocolPB.ClientNamenodeProtocolServerSideTranslatorPB.addBlock(ClientNamenodeProtocolServerSideTr
anslatorPB.java:603)
    at org.apache.hadoop.hdfs.protocol.proto.ClientNamenodeProtocolProtos$ClientNamenodeProtocol$2.callBlockingMethod(ClientNamen
odeProtocolProtos.java)
    at org.apache.hadoop.ipc.ProtobufRpcEngine2$Server$ProtoBufRpcInvoker.call(ProtobufRpcEngine2.java:621)
    at org.apache.hadoop.ipc.ProtobufRpcEngine2$Server$ProtoBufRpcInvoker.call(ProtobufRpcEngine2.java:589)
    at org.apache.hadoop.ipc.ProtobufRpcEngine2$Server$ProtoBufRpcInvoker.call(ProtobufRpcEngine2.java:573)
    at org.apache.hadoop.ipc.RPC$Server.call(RPC.java:1227)
    at org.apache.hadoop.ipc.Server$RpcCall.run(Server.java:1246)
    at org.apache.hadoop.ipc.Server$RpcCall.run(Server.java:1169)
    at java.security.AccessController.doPrivileged(Native Method)
    at javax.security.auth.Subject.doAs(Subject.java:422)
    at org.apache.hadoop.security.UserGroupInformation.doAs(UserGroupInformation.java:1953)
    at org.apache.hadoop.ipc.Server$Handler.run(Server.java:3203)

    at org.apache.hadoop.ipc.Client.getRpcResponse(Client.java:1584)

```

```
2024-09-30 19:56:33,780 WARN datanode.DataNode: Ending block pool serv
4916-bf64-e5792c2b97be) service to localhost/127.0.0.1:9000
2024-09-30 19:56:33,781 INFO datanode.DataNode: Removed Block pool <re
2024-09-30 19:56:35,781 WARN datanode.DataNode: Exiting Datanode
2024-09-30 19:56:35,791 INFO datanode.DataNode: SHUTDOWN_MSG:
/*****
SHUTDOWN_MSG: Shutting down DataNode at LAPTOP-UGU6DF82/127.0.1.1
*****/
deepsalunkhee@LAPTOP-UGU6DF82:~$ $ stop-yarn.sh
$ stop-dfs.sh
$: command not found
$: command not found
deepsalunkhee@LAPTOP-UGU6DF82:~$ stop-dfs.sh
Stopping namenodes on [localhost]
Stopping datanodes
Stopping secondary namenodes [LAPTOP-UGU6DF82]
deepsalunkhee@LAPTOP-UGU6DF82:~$ stop-yarn.sh
Stopping nodemanagers
Stopping resourcemanager
deepsalunkhee@LAPTOP-UGU6DF82:~$ █
```