

- Name: Deep Salunkhe
- Roll No.:21102A0014
- [SEM-7 ML Lab5 Github Link](#)

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.svm import SVC, SVR
from sklearn.metrics import accuracy_score, precision_score,
recall_score, f1_score, mean_squared_error, r2_score
from sklearn.metrics import confusion_matrix, classification_report

# Load datasets
attributes_array = [
    'fixed acidity', 'volatile acidity', 'citric acid',
    'residual sugar', 'chlorides', 'free sulfur dioxide',
    'total sulfur dioxide', 'density', 'pH', 'sulphates',
    'alcohol', 'quality'
]

df_red = pd.read_csv('/content/winequality-red.csv', delimiter=";")
df_white = pd.read_csv('/content/winequality-white.csv',
delimeter=";")

df = pd.concat([df_red, df_white])

df.head()

{"summary":{"\n  \"name\": \"df\", \n  \"rows\": 6497, \n  \"fields\":
[\n    {\n      \"column\": \"fixed acidity\", \n      \"properties\":
{\n        \"dtype\": \"number\", \n        \"std\":
1.2964337577998153, \n        \"min\": 3.8, \n        \"max\": 15.9, \n
\n      \"num_unique_values\": 106, \n      \"samples\": [\n        7.15, \n
8.1, \n        7.3 \n      ], \n      \"semantic_type\": \"\", \n
\n      \"description\": \"\" \n    }, \n    {\n      \"column\":
\"volatile acidity\", \n      \"properties\": {\n        \"dtype\":
\"number\", \n        \"std\": 0.16463647408467877, \n        \"min\":
0.08, \n        \"max\": 1.58, \n        \"num_unique_values\": 187, \n
\n      \"samples\": [\n        0.405, \n        0.21, \n        0.695 \n
      ], \n      \"semantic_type\": \"\", \n      \"description\": \"\" \n
    }, \n    {\n      \"column\": \"citric acid\", \n      \"properties\": {\n
        \"dtype\": \"number\", \n        \"std\":
0.14531786489759155, \n        \"min\": 0.0, \n        \"max\": 1.66, \n
\n      \"num_unique_values\": 89, \n      \"samples\": [\n        0.1, \n
0.6, \n        0.37 \n      ], \n      \"semantic_type\": \"\", \n
\n    } \n  ] \n  } \n}
```

```

{"description": "", "properties": {"column": "residual sugar", "dtype": "number", "std": 4.757803743147418, "min": 0.6, "max": 65.8, "num_unique_values": 316, "samples": [18.95, 3.2, 9.3]}, "semantic_type": "", "description": ""}, {"column": "chlorides", "dtype": "number", "std": 0.03503360137245907, "min": 0.009, "max": 0.611, "num_unique_values": 214, "samples": [0.089, 0.217, 0.1]}, "semantic_type": "", "description": ""}, {"column": "free sulfur dioxide", "dtype": "number", "std": 17.7493997720025, "min": 1.0, "max": 289.0, "num_unique_values": 135, "samples": [77.5, 65.0, 128.0]}, "semantic_type": "", "description": ""}, {"column": "total sulfur dioxide", "dtype": "number", "std": 56.521854522630285, "min": 6.0, "max": 440.0, "num_unique_values": 276, "samples": [14.0, 149.0, 227.0]}, "semantic_type": "", "description": ""}, {"column": "density", "dtype": "number", "std": 0.0029986730037190393, "min": 0.98711, "max": 1.03898, "num_unique_values": 998, "samples": [0.9918, 0.99412, 0.99484]}, "semantic_type": "", "description": ""}, {"column": "pH", "dtype": "number", "std": 0.16078720210398764, "min": 2.72, "max": 4.01, "num_unique_values": 108, "samples": [3.74, 3.17, 3.3]}, "semantic_type": "", "description": ""}, {"column": "sulphates", "dtype": "number", "std": 0.14880587361449027, "min": 0.22, "max": 2.0, "num_unique_values": 111, "samples": [1.11, 1.56, 0.46]}, "semantic_type": "", "description": ""}, {"column": "alcohol", "dtype": "number", "std": 1.192711748870993, "min": 8.0, "max": 14.9, "num_unique_values": 111, "samples": [9.7, 10.5]}, "semantic_type": "", "description": ""}, {"column": "quality", "dtype": "number", "std": 0, "min": 3, "max": 9, "num_unique_values": 7, "samples": []}, {"semantic_type": "", "description": ""}

```

```

\"samples\": [\n                5,\n                6,\n                3\n            ],\n\"semantic_type\": \"\",\n\"description\": \"\"\n    }\n    ],\n    },\n    \"type\": \"dataframe\", \"variable_name\": \"df\"}

df['quality'].unique()
array([5, 6, 7, 4, 8, 3, 9])

df.isnull().sum()
fixed acidity      0
volatile acidity   0
citric acid        0
residual sugar     0
chlorides          0
free sulfur dioxide 0
total sulfur dioxide 0
density            0
pH                0
sulphates          0
alcohol            0
quality            0
dtype: int64


bins = (2, 6.5, 9)
group_names = ['bad', 'good']
df['quality'] = pd.cut(df['quality'], bins = bins, labels =
group_names)
df['quality'].unique()

['bad', 'good']
Categories (2, object): ['bad' < 'good']

label_quality = LabelEncoder()
df['quality'] = label_quality.fit_transform(df['quality'])

sns.countplot(x='quality', data=df)
plt.show()

```



```
#Now separate the dataset as response variable and feature variables
X = df.drop('quality', axis=1)
y = df['quality']

#Train and test splitting of data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size =
0.2, random_state = 42)

#Applying Standard scaling to get optimized result
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)

# Create an SVM model
svm_model = SVC(kernel='linear', C=1)

# Train the model
svm_model.fit(X_train, y_train)

# Make predictions on the test set
y_pred = svm_model.predict(X_test)

# Calculate evaluation metrics
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred, average = 'weighted',
```

```

zero_division = 1)
recall = recall_score(y_test, y_pred, average = 'weighted')
f1 = f1_score(y_test, y_pred, average = 'weighted')

# Print the results
print("Accuracy:", accuracy)
print("Precision:", precision)
print("Recall:", recall)
print("F1-score:", f1)

Accuracy: 0.8061538461538461
Precision: 0.843730177514793
Recall: 0.8061538461538461
F1-score: 0.7196330756126327

print(classification_report(y_test, y_pred, zero_division=1))
print(confusion_matrix(y_test, y_pred))

```

	precision	recall	f1-score	support
0	0.81	1.00	0.89	1048
1	1.00	0.00	0.00	252
accuracy			0.81	1300
macro avg	0.90	0.50	0.45	1300
weighted avg	0.84	0.81	0.72	1300

```

[[1048    0]
 [ 252    0]]

from sklearn.preprocessing import label_binarize

# Binarize the output labels for ROC curve calculation
y_test_bin = label_binarize(y_test, classes=[0, 1])
y_pred_prob = svm_model.decision_function(X_test)

# Compute ROC curve and AUC for each class
fpr = dict()
tpr = dict()
roc_auc = dict()


fpr, tpr, _ = roc_curve(y_test_bin, y_pred_prob)
roc_auc = auc(fpr, tpr)

# Plot ROC curve
plt.figure()
lw = 2
plt.plot(fpr, tpr, color='darkorange',
         lw=lw, label='ROC curve (area = %0.2f)' % roc_auc)
plt.plot([0, 1], [0, 1], color='navy', lw=lw, linestyle='--')
plt.xlim([0.0, 1.0])

```

```
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic')
plt.legend(loc="lower right")
plt.show()

print("AUC:", roc_auc)
```



AUC: 0.684326911426148

## Regression

```
# Load the dataset
df_concrete = pd.read_excel('/content/Concrete_Data.xls')

attributes_array = [
    'Cement',
    'Blast Furnace Slag',
    'Fly Ash',
    'Water',
```

```

        'Superplasticizer',
        'Coarse Aggregate',
        'Fine Aggregate',
        'Age',
        'Concrete compressive strength'
    ]

df_concrete.columns = attributes_array
df_concrete.to_csv('concrete_data.csv', index=False)
df = pd.read_csv('concrete_data.csv')
df.head()

{"summary": "{\n  \"name\": \"df\",\n  \"rows\": 1030,\n  \"fields\": [\n    {\n      \"column\": \"Cement\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 104.5071416428718,\n        \"min\": 102.0,\n        \"max\": 540.0,\n        \"num_unique_values\": 280,\n        \"samples\": [\n          194.68,\n          480.0,\n          145.4\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"Blast Furnace Slag\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 86.27910364316895,\n        \"min\": 0.0,\n        \"max\": 359.4,\n        \"num_unique_values\": 187,\n        \"samples\": [\n          186.7,\n          212.0,\n          26.0\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"Fly Ash\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 63.99646938186508,\n        \"min\": 0.0,\n        \"max\": 200.1,\n        \"num_unique_values\": 163,\n        \"samples\": [\n          81.8,\n          137.9,\n          107.5\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"Water\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 21.355567066911522,\n        \"min\": 121.75,\n        \"max\": 247.0,\n        \"num_unique_values\": 205,\n        \"samples\": [\n          164.9,\n          181.1,\n          185.7\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"Superplasticizer\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 5.973491650590111,\n        \"min\": 0.0,\n        \"max\": 32.2,\n        \"num_unique_values\": 155,\n        \"samples\": [\n          4.14,\n          9.8,\n          6.13\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"Coarse Aggregate\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 77.75381809178927,\n        \"min\": 801.0,\n        \"max\": 1145.0,\n        \"num_unique_values\": 284,\n        \"samples\": [\n          852.1,\n          913.9,\n          914.0\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"Fine Aggregate\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\":

```

```

0.1754273990239,\n                \"min\": 594.0,\n                \"max\": 992.6,\n                \"num_unique_values\": 304,\n                \"samples\": [\n                    613.0,\n                    689.3\n                ],\n                \"semantic_type\": \"\",\n                \"description\": \"\",\n                \"column\": \"Age\",\n                \"properties\": {\n                    \"dtype\": \"number\",\n                    \"std\": 63,\n                    \"min\": 1,\n                    \"max\": 365,\n                    \"num_unique_values\": 14,\n                    \"samples\": [\n                        91,\n                        100,\n                        28\n                    ],\n                    \"semantic_type\": \"\",\n                    \"description\": \"\",\n                    \"column\": \"Concrete compressive strength\",\n                    \"properties\": {\n                        \"dtype\": \"number\",\n                        \"std\": 16.705679174867946,\n                        \"min\": 2.331807832,\n                        \"max\": 82.5992248,\n                        \"num_unique_values\": 938,\n                        \"samples\": [\n                            33.39821744,\n                            56.63355864,\n                            25.559564796\n                        ],\n                        \"semantic_type\": \"\",\n                        \"description\": \"\"\n                    }\n                }\n            ],\n            \"type\": \"dataframe\",\n            \"variable_name\": \"df\"\n        }
    ]
}

```

```
# Increase the figure size for better readability
plt.figure(figsize=(12, 10))
```

```
# Create the heatmap with annotations and a title
sns.heatmap(df.corr(), annot=True, linewidths=.5, fmt=".2f",
cmap="coolwarm")
plt.title("Correlation Matrix of Concrete Data")
plt.show()
```


```
print()
```

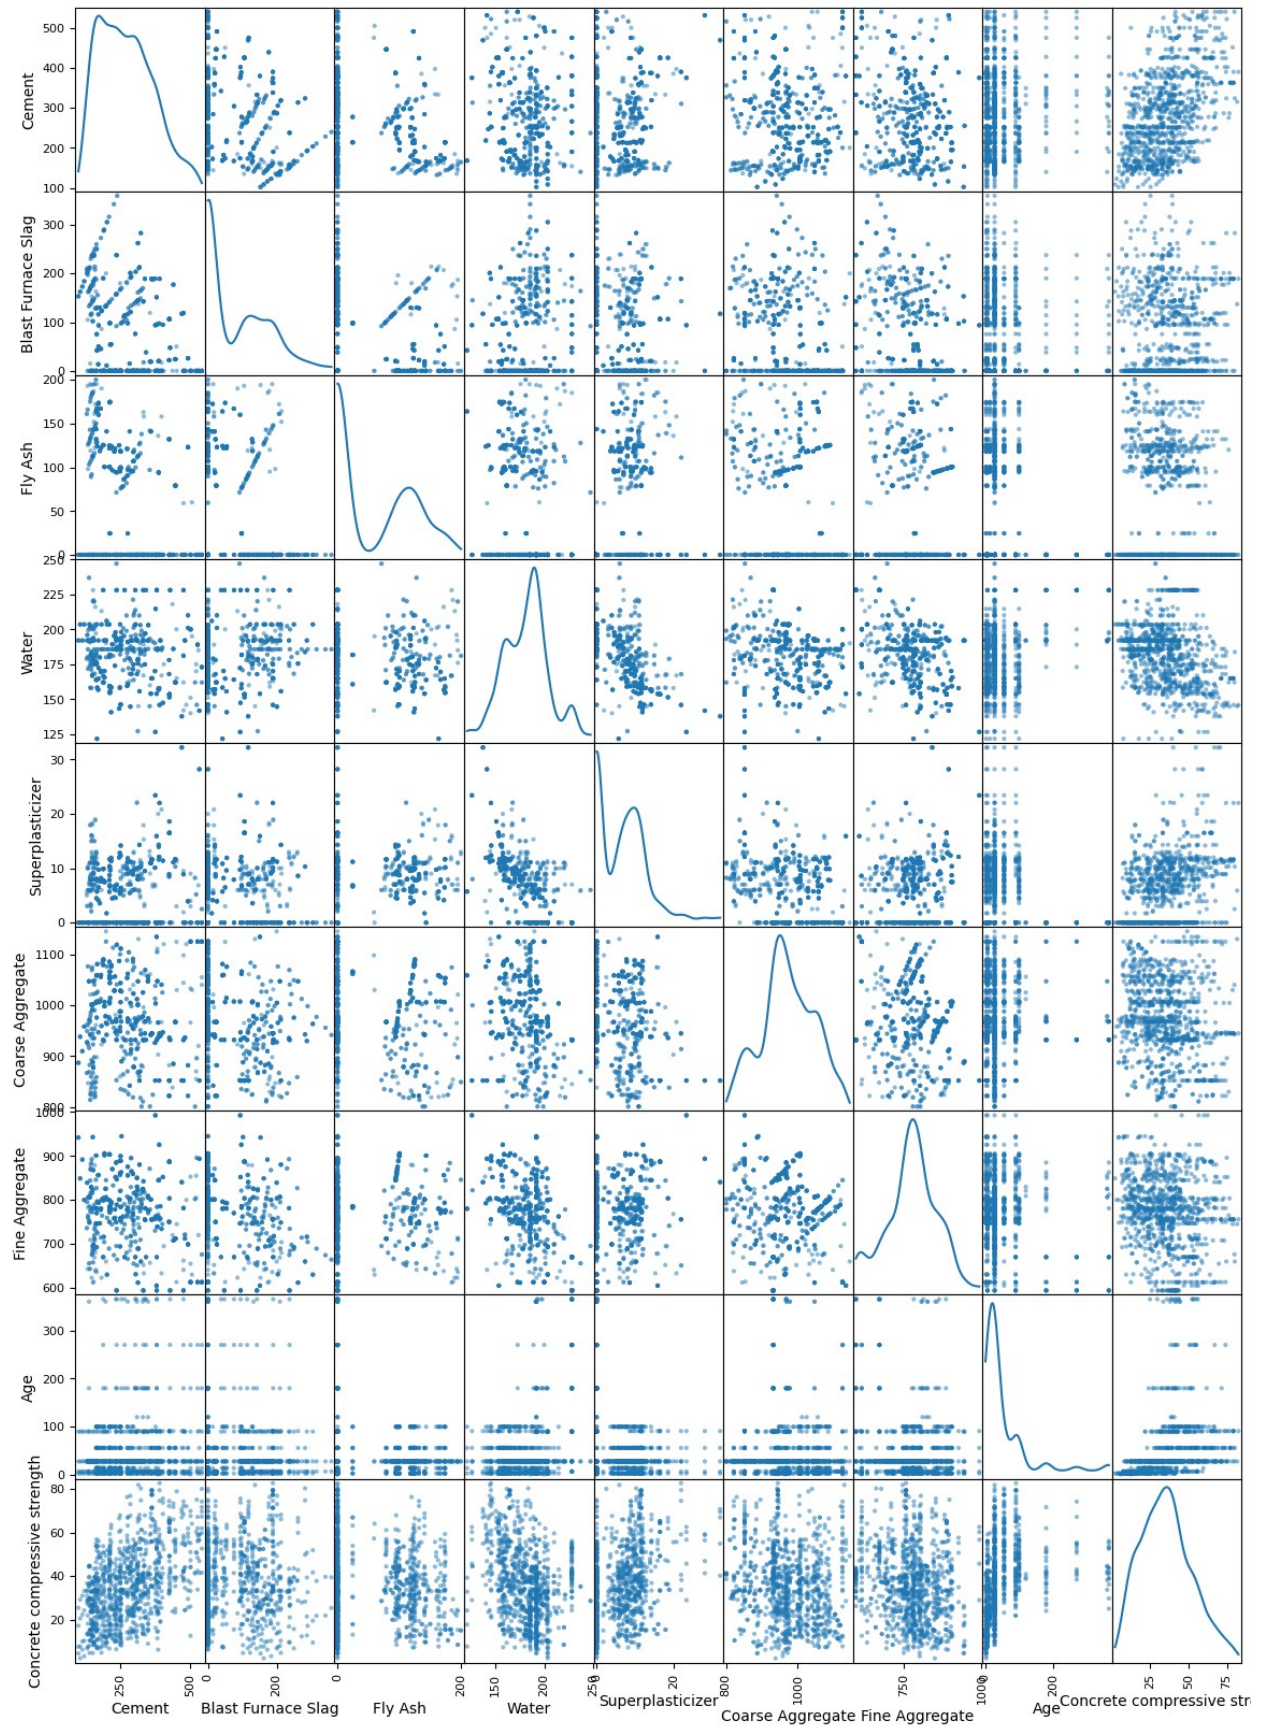
```
from pandas.plotting import scatter_matrix
scatter_matrix(df, figsize=(14, 20), diagonal='kde')
plt.show()
```


```
print()
```

```
# Create a distribution plot with increased figure size and a
# descriptive title
plt.figure(figsize=(10, 6))
sns.histplot(df['Concrete compressive strength'], bins=15, kde=True,
color='skyblue')
plt.xlabel("Concrete Compressive Strength (MPa)")
plt.ylabel("Frequency")
plt.title('Distribution of Concrete Compressive Strength')
plt.show()
```









```
# Select columns for visualization (excluding 'Concrete compressive strength')
cols = [col for col in df.columns if col != 'Concrete compressive strength']




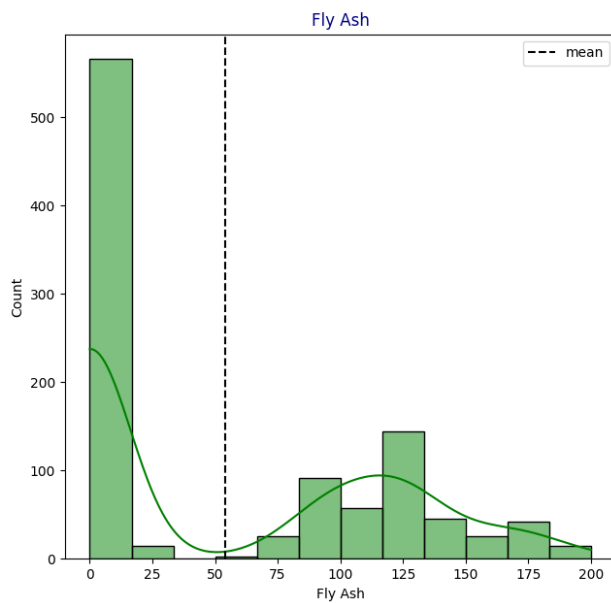
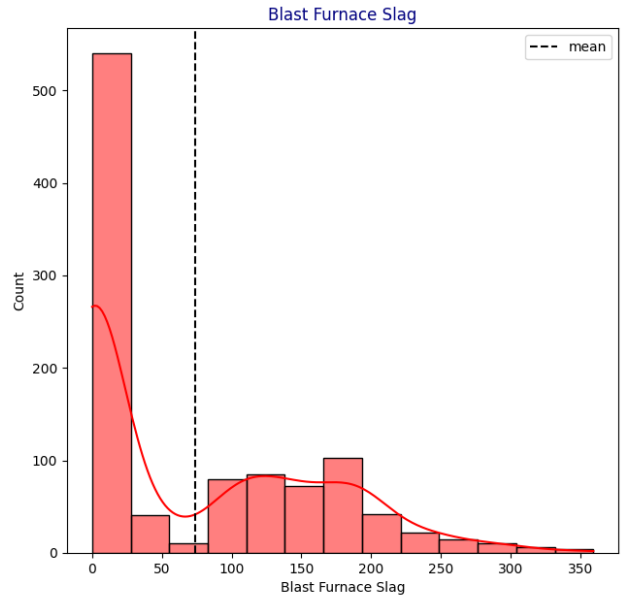
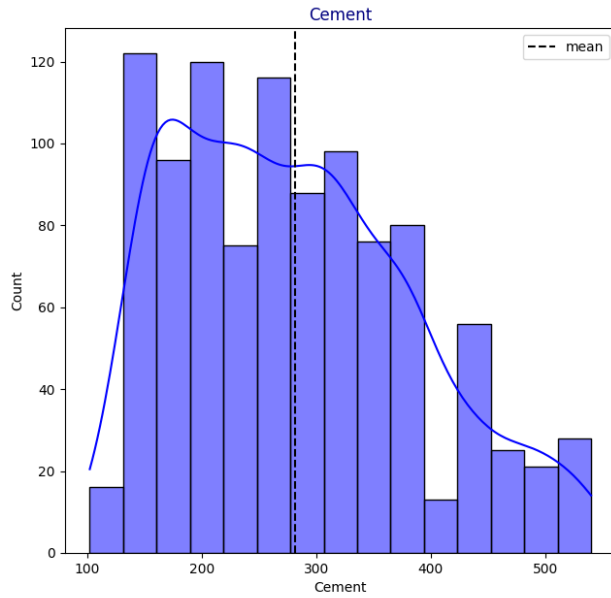
# Create a figure with subplots
fig, axes = plt.subplots(nrows=4, ncols=2, figsize=(13, 25))
axes = axes.flatten()

# Define colors for plots
colors = ['b', 'r', 'g', 'c', 'm', 'k', 'lime', 'c']

# Iterate through columns and create distribution plots
for i, col in enumerate(cols):
    sns.histplot(df[col], color=colors[i], kde=True, ax=axes[i])
    axes[i].set_facecolor("w")
    axes[i].axvline(df[col].mean(), linestyle="dashed", label="mean", color="k")
    axes[i].set_title(col, color="navy")
    axes[i].legend()

# Adjust layout and show plot
```

```
plt.tight_layout()  
plt.show()
```



```

# Split the data into features (X) and target (y)
X_concrete = df.drop('Concrete compressive strength', axis=1)
y_concrete = df['Concrete compressive strength']

# Split the data into training and testing sets
X_train_concrete, X_test_concrete, y_train_concrete, y_test_concrete =
train_test_split(
    X_concrete, y_concrete, test_size=0.3, random_state=0
)

# Scale the features using StandardScaler
scaler_concrete = StandardScaler()
X_train_concrete = scaler_concrete.fit_transform(X_train_concrete)
X_test_concrete = scaler_concrete.transform(X_test_concrete)

# Create an SVR model
svr_model = SVR(kernel='linear', C=1)

# Train the model
svr_model.fit(X_train_concrete, y_train_concrete)

# Make predictions on the test set
y_pred_concrete = svr_model.predict(X_test_concrete)

# Evaluate the model
mse = mean_squared_error(y_test_concrete, y_pred_concrete)
r2 = r2_score(y_test_concrete, y_pred_concrete)

print("SVR Model Results:")
print("Mean Squared Error:", mse)
print("R-squared:", r2)

SVR Model Results:
Mean Squared Error: 93.07450744225504
R-squared: 0.6374191118739791

from sklearn.preprocessing import PolynomialFeatures

poly = PolynomialFeatures(degree=3) # Try different degrees
X_train_concrete_poly = poly.fit_transform(X_train_concrete)
X_test_concrete_poly = poly.transform(X_test_concrete)

svr_model.fit(X_train_concrete_poly, y_train_concrete)
y_pred_concrete = svr_model.predict(X_test_concrete_poly)

# Evaluate the model
mse = mean_squared_error(y_test_concrete, y_pred_concrete)
r2 = r2_score(y_test_concrete, y_pred_concrete)

print("SVR Model Results:")

```

```
print("Mean Squared Error:", mse)
print("R-squared:", r2)

SVR Model Results:
Mean Squared Error: 44.55026320014946
R-squared: 0.8264500727293151

# Create a DataFrame from the actual and predicted values
dat = pd.DataFrame({'Actual': y_test_concrete, 'Predicted':
y_pred_concrete})

# Sort the DataFrame by the 'Actual' column
dat_sorted = dat.sort_values(by=['Actual'])

# Scatter plot of Actual vs Predicted
plt.figure(figsize=(7, 7))

plt.scatter(dat_sorted['Actual'], dat_sorted['Predicted'],
color='blue', label='Predicted')

# Plot the perfect prediction line (Actual = Predicted)
plt.plot(dat_sorted['Actual'], dat_sorted['Actual'], color='red',
label='Perfect Prediction', linestyle='--')

# Add grid, labels, and title
plt.grid(which='major', linestyle='-', linewidth='0.5', color='green')
plt.title('Actual vs Predicted (Scatter Plot)')
plt.xlabel('Actual Values')
plt.ylabel('Predicted Values')

# Add a legend
plt.legend()

# Show the plot
plt.show()
```



Actual vs Predicted (Scatter Plot)

