

**Assignment No. 11**

Semester	B.E. Semester VIII – Computer Engineering
Subject	Distributed Computing Lab
Subject Professor In-charge	Dr. Umesh Kulkarni
Assisting Professor	Prof. Prakash Parmar
Academic Year	2024-25
Student Name	Deep Salunkhe
Roll Number	21102A0014

**Title:** Balancing Trade-offs in Distributed File Systems (DFS)

---

Distributed File Systems (DFS) aim to provide efficient, reliable, and scalable access to files over a network, but achieving this requires balancing **efficient file access, caching, replication, and fault tolerance**. Modern DFSs, such as **NFS (Network File System)** and **Google File System (GFS)**, implement various strategies to optimize performance while ensuring data consistency and availability.

---

## 1. Trade-offs in Distributed File Systems (DFS)

DFSs must balance the following key factors:

### A. Efficient File Access

DFSs aim to minimize latency and maximize throughput when retrieving or writing files.

- **Challenges:**
  - Network delays can slow down file access.
  - Centralized metadata management may become a bottleneck.
- **Solutions:**

- **Chunk-based storage (GFS, HDFS):** Large files are split into chunks, allowing parallel access.
- **Distributed caching (NFS, AFS):** Frequently accessed files are cached at the client side.
- **Optimized read/write mechanisms:** Pre-fetching and asynchronous writes improve performance.

**Example: Google File System (GFS)** divides files into 64MB chunks stored on multiple chunk servers, reducing metadata overhead and improving parallel access.

---

## B. Caching Mechanisms

Caching helps improve performance by reducing redundant data fetches from remote storage.

- **Challenges:**

- **Cache consistency:** How to ensure updates from different clients remain consistent?
- **Stale data:** Cached copies may become outdated if the original file changes.

- **Solutions:**

- **Write-through caching (NFSv3):** Writes are immediately sent to the server, ensuring consistency.
- **Client-side caching (AFS, NFSv4):** Local copies reduce access latency.
- **Lease-based caching (NFSv4, GFS):** The server grants leases to clients, allowing temporary exclusive access.

**Example: NFSv4 uses delegation,** where the server temporarily gives clients the authority to cache files, reducing network load.

---

## C. Replication for Fault Tolerance & Availability

Replication ensures data remains available even if some nodes fail.

- **Challenges:**
  - **Replication overhead:** Maintaining multiple copies requires extra storage and synchronization.
  - **Consistency issues:** Updates must be propagated to all replicas correctly.
- **Solutions:**
  - **Primary-backup replication (GFS, HDFS):** A primary copy manages writes, and updates are propagated to secondary replicas.
  - **Eventual consistency (Amazon S3, Cassandra):** Writes are asynchronously replicated to ensure availability.
  - **Quorum-based replication (Ceph, Google Spanner):** Ensures strong consistency by requiring a majority of replicas to acknowledge changes.

**Example: Google File System (GFS)** replicates chunks **three times** across different servers for fault tolerance.

---

## D. Fault Tolerance & Recovery

DFSs must ensure availability despite node failures.

- **Challenges:**
  - **Detecting and recovering from failures efficiently.**
  - **Avoiding single points of failure.**
- **Solutions:**
  - **Metadata replication (GFS, Ceph):** Ensures the metadata server is not a single point of failure.
  - **Heartbeats & automatic failover (HDFS, GFS):** Regular health checks allow quick failover to backup nodes.
  - **Erase coding (Ceph, HDFS):** Reduces storage overhead while providing fault tolerance.

**Example: Google's GFS master node regularly checkpoints metadata** and assigns

a new master in case of failure.

---

## 2. How Modern DFSs Implement These Trade-offs

Feature	NFS	Google File System (GFS)
<b>Efficient File Access</b>	Client-server architecture, optimized for local networks	Chunk-based file storage with parallel access
<b>Caching</b>	Client-side caching with consistency mechanisms	Lease-based caching for efficient reads
<b>Replication</b>	External replication using RAID or backups	Three-way chunk replication for fault tolerance
<b>Fault Tolerance</b>	Stateless server (NFSv3), Stateful in (NFSv4)	Master node failure detection and automatic recovery
<b>Scalability</b>	Best for enterprise LANs	Optimized for massive-scale web applications

---

## 3. Conclusion: Achieving a Balanced DFS Design

A well-designed **Distributed File System** must optimize **efficiency, caching, replication, and fault tolerance** while adapting to application needs.

- **NFS is optimized for network file sharing**, ensuring consistency and security.
- **GFS prioritizes scalability and fault tolerance** with chunk-based storage and replication.
- **Modern DFSs (HDFS, Ceph, Amazon S3) continue evolving** by integrating cloud-based storage, machine learning-based optimizations, and erasure coding.